

Anchored Diffusion Language Model

NeurIPS 2025

Litu Rout, Constantine Caramanis, and Sanjay Shakkottai

The University of Texas at Austin



The University of Texas at Austin
Machine Learning Laboratory



The University of Texas at Austin
Chandra Department of Electrical
and Computer Engineering
Cockrell School of Engineering

IFML

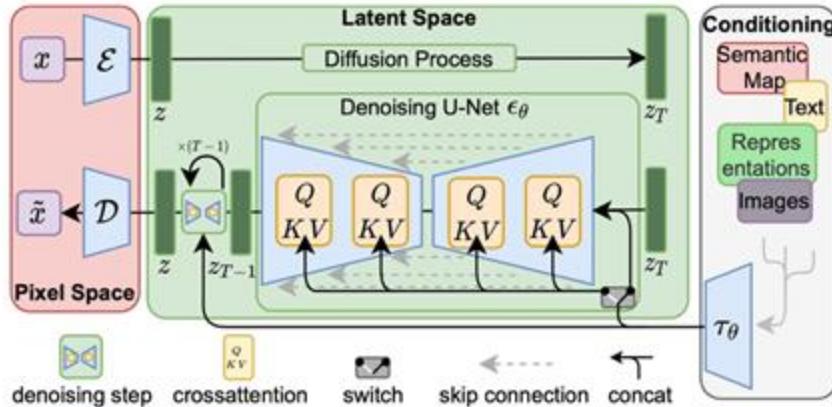
AI EDGE
Institute

LDOS

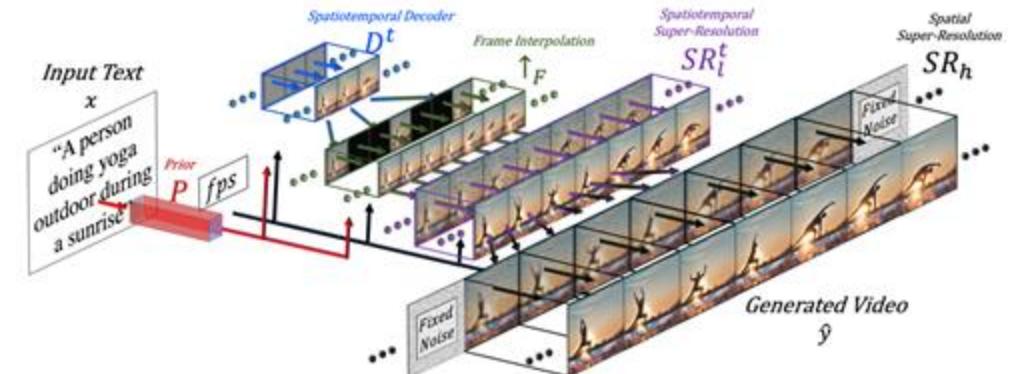


Generative Modeling with Diffusion

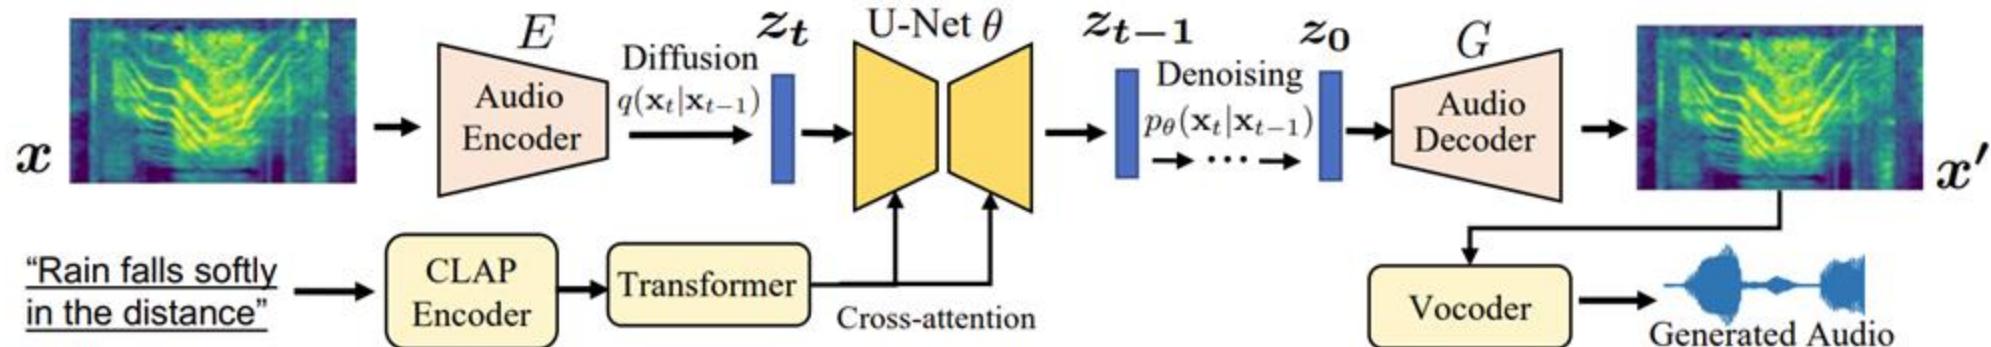
Diffusion Model for Images



Diffusion Model for Video



Diffusion Model for Audio

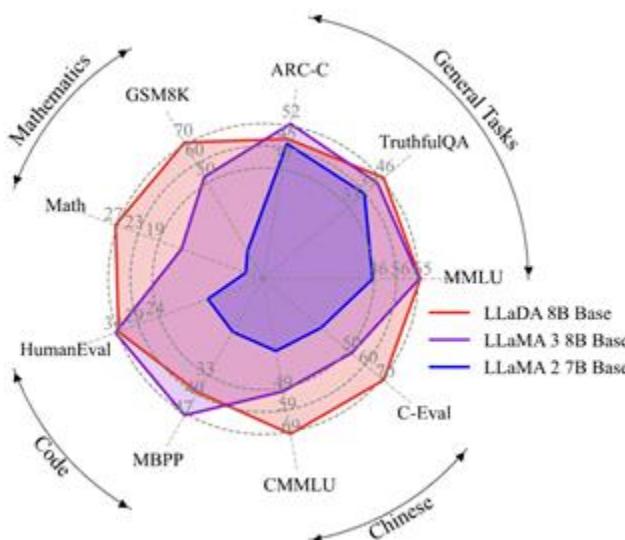


Large-Scale Diffusion Language Models

Gemini Diffusion



LLaDA



Mercury

Can be more accurate

Significantly faster

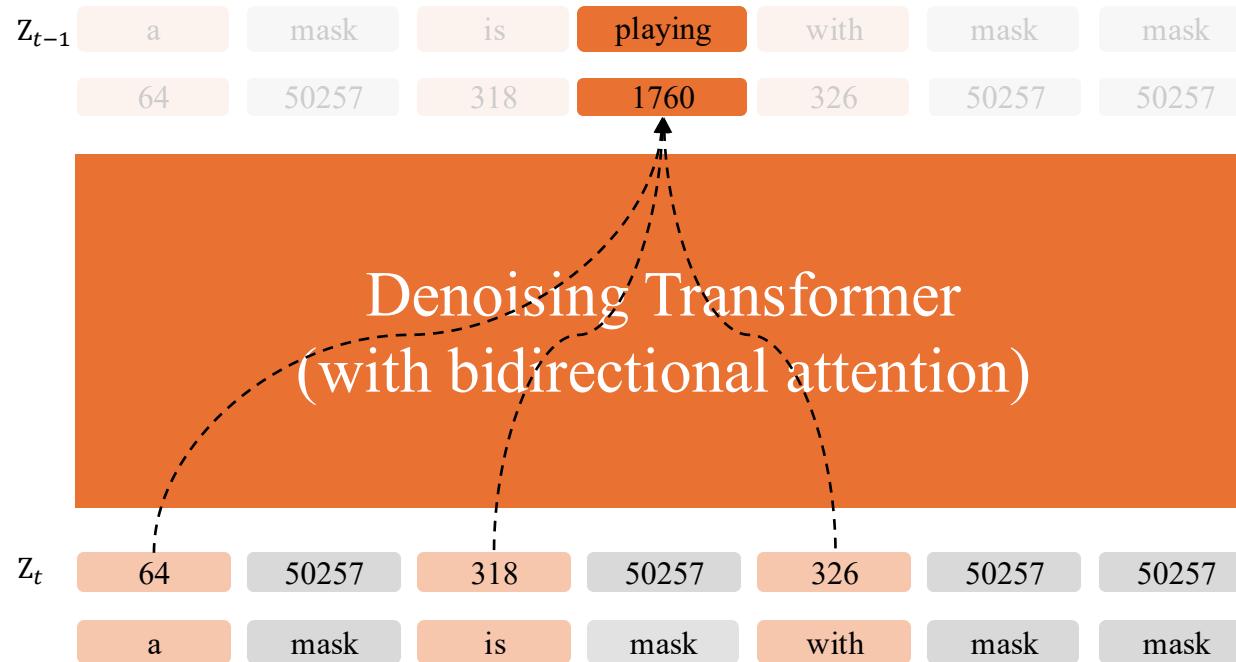
The image compares the training times for Mercury and LLaDA. On the left, the LLaDA training process is shown with 75 iterations and the text "AUTOREGRESSIVE LLM LEFT-TO-RIGHT GENERATION". On the right, the Mercury training process is shown with 14 iterations and the text "INCEPTION DIFFUSION LLM COARSE-TO-FINE GENERATION". Both processes are labeled as "Completed". Above the panels, a callout box says "Write a function for LLM inference." Below the panels, a large callout box at the top right says "DIFFUSION IS FAR FASTER.".

```
def auto_regressive_decode(model, start_token, max_length):
    sequence = [start_token]
    while len(sequence) < max_length:
        input_seq = torch.tensor([sequence])
        output = model(input_seq)
        next_token = torch.argmax(output, dim=-1).item()
        sequence.append(next_token)
        if next_token == model.eos_token_id:
            break
    return sequence
```

```
def auto_regressive_decode(model, start_token, max_length):
    sequence = [start_token]
    while len(sequence) < max_length:
        input_seq = torch.tensor([sequence])
        output = model(input_seq)
        next_token = torch.argmax(output, dim=-1).item()
        sequence.append(next_token)
        if next_token == model.eos_token_id:
            break
    return sequence
```

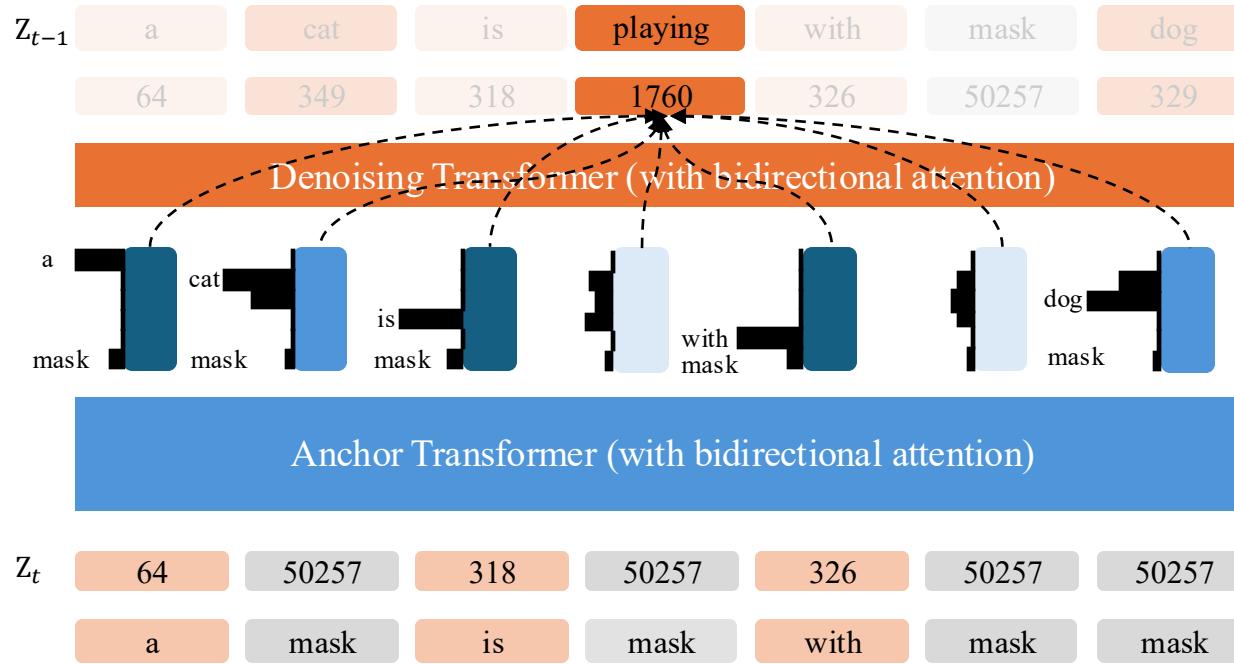
Gemini Diffusion: <https://deepmind.google/models/gemini-diffusion/>
LLaDA: <https://arxiv.org/abs/2502.09992>
Mercury: <https://www.inceptionlabs.ai/>

Standard Approach: Masked Diffusion Language Model



- **Denoising Transformer:** Unmasks using already unmasked tokens
- **Our approach:** Improve context using semantically important tokens

Our Approach: Anchored Diffusion Language Model

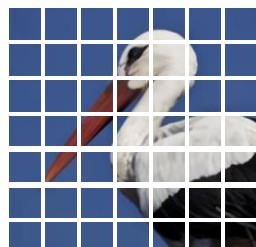


- **Anchor Transformer:** Outputs a sequence of anchor predictions
 - Mixture of important tokens interpreted as soft samples
- **Denoising Transformer:** Unmasks tokens using anchored predictions

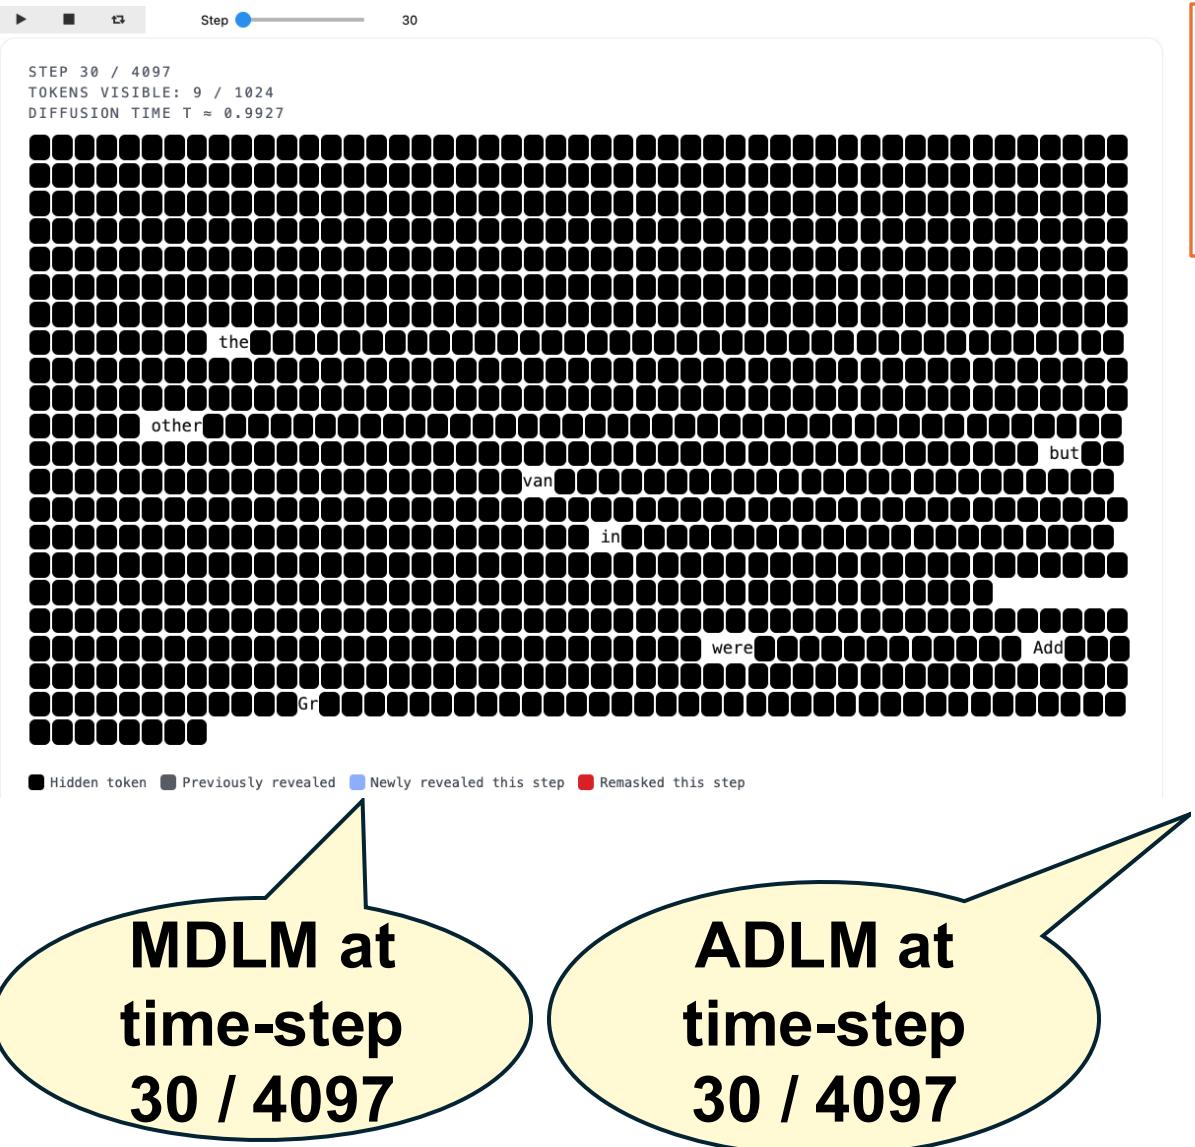
Our Key Idea: Anchoring

Anchors: Tokens whose inclusion as conditioning variables yields a **substantial reduction** in the **conditional entropy** of the remaining tokens

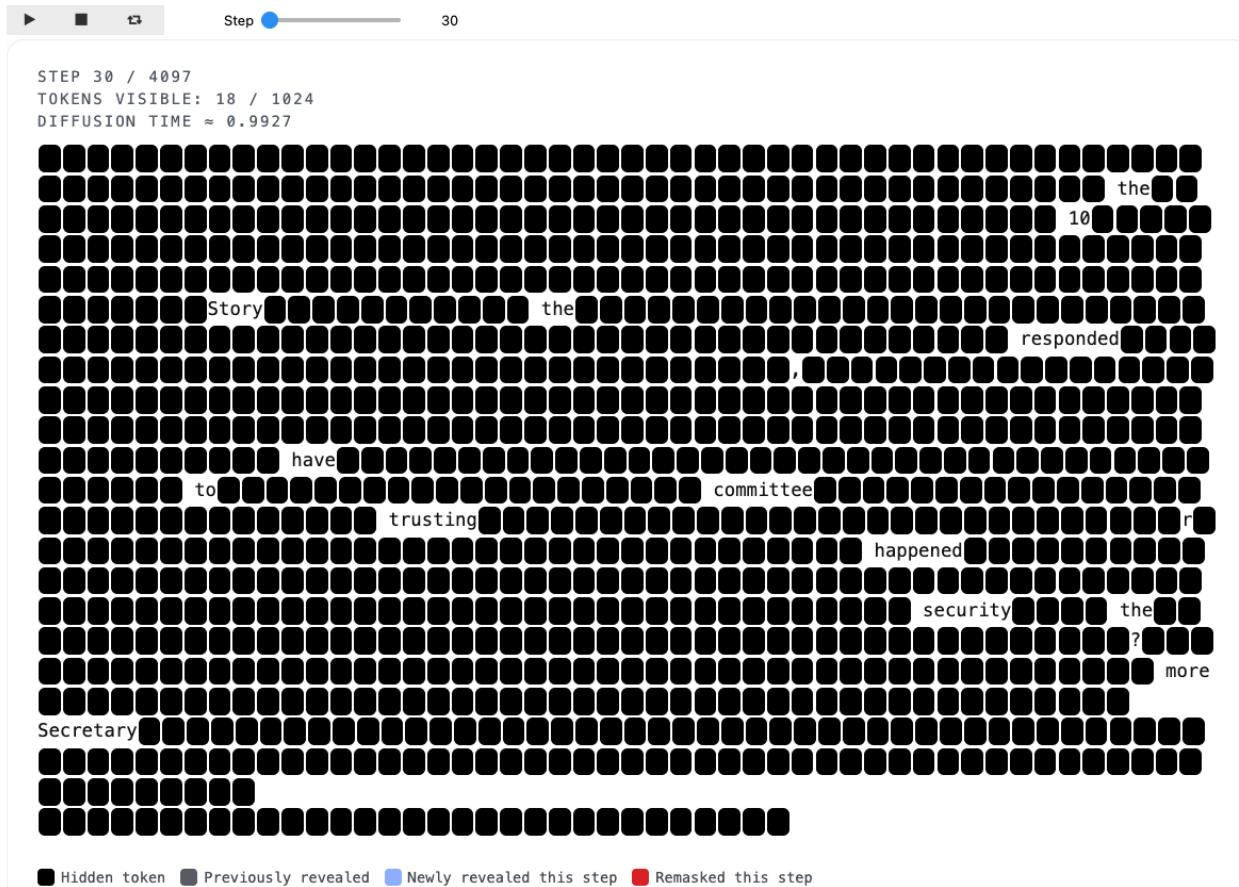
Examples:

1. To decode the sentence: “a **cat** is playing with a **dog**”
 - Tokens ‘**cat**’ and ‘**dog**’ are more useful than ‘**a**’ or ‘**is**’
2. To solve the math question: “Janet’s ducks lay **16** eggs per day. She eats **three** for breakfast every morning and bakes muffins for her friends every day with **four**. She sells the remainder at the farmers’ market daily for **\$2** per fresh duck egg. How much in dollars does she make every day at the farmers’ market?”
 - Numbers ‘**16**’, ‘**three**’, ‘**four**’, and ‘**2**’ are more useful than ‘**breakfast**’ or ‘**muffins**’
3. To reconstruct the image:
 - Tokens in the **beak** or **body** (foreground) are more useful than tokens from blue **background**

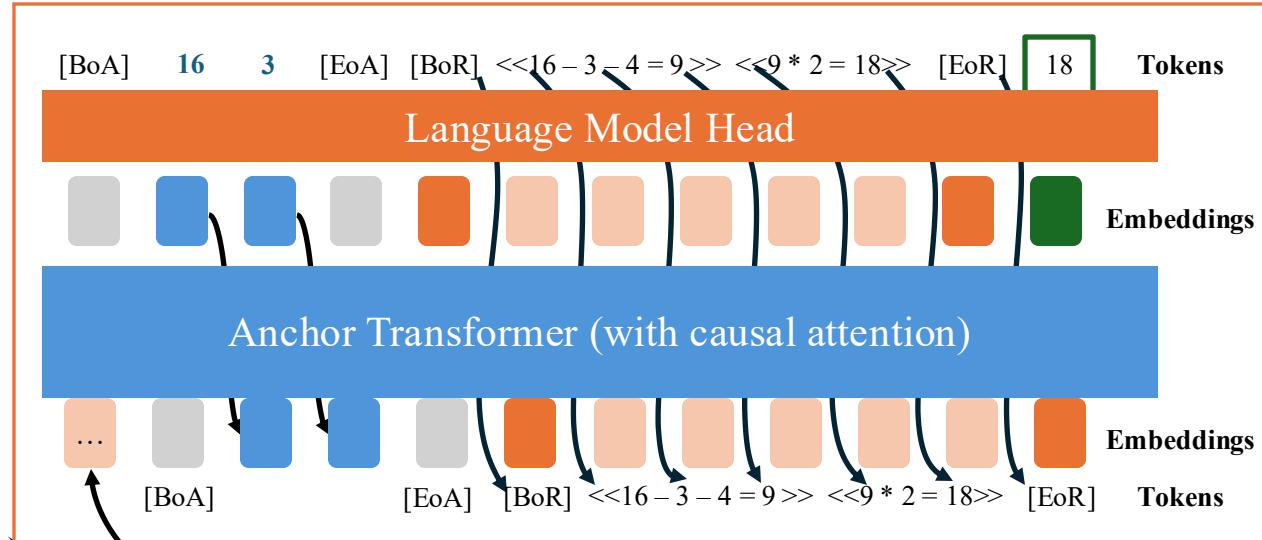
Masked DLM vs Anchored DLM – Inference Illustration



- Anchoring has 2 main advantages:
 - unmasks **key words** first
 - unmasks **many tokens** in parallel



Anchored Autoregressive Model: Training and Inference

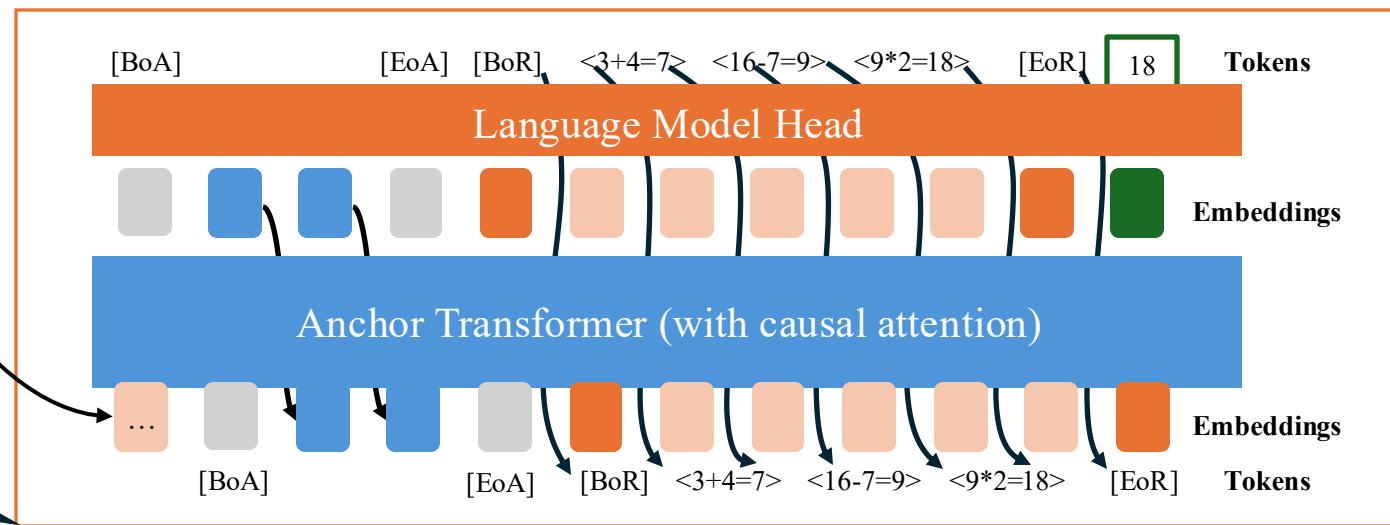


Anchoring enables “look-ahead” planning, unlike the standard left-to-right decoding of chain-of-thought (CoT) reasoning

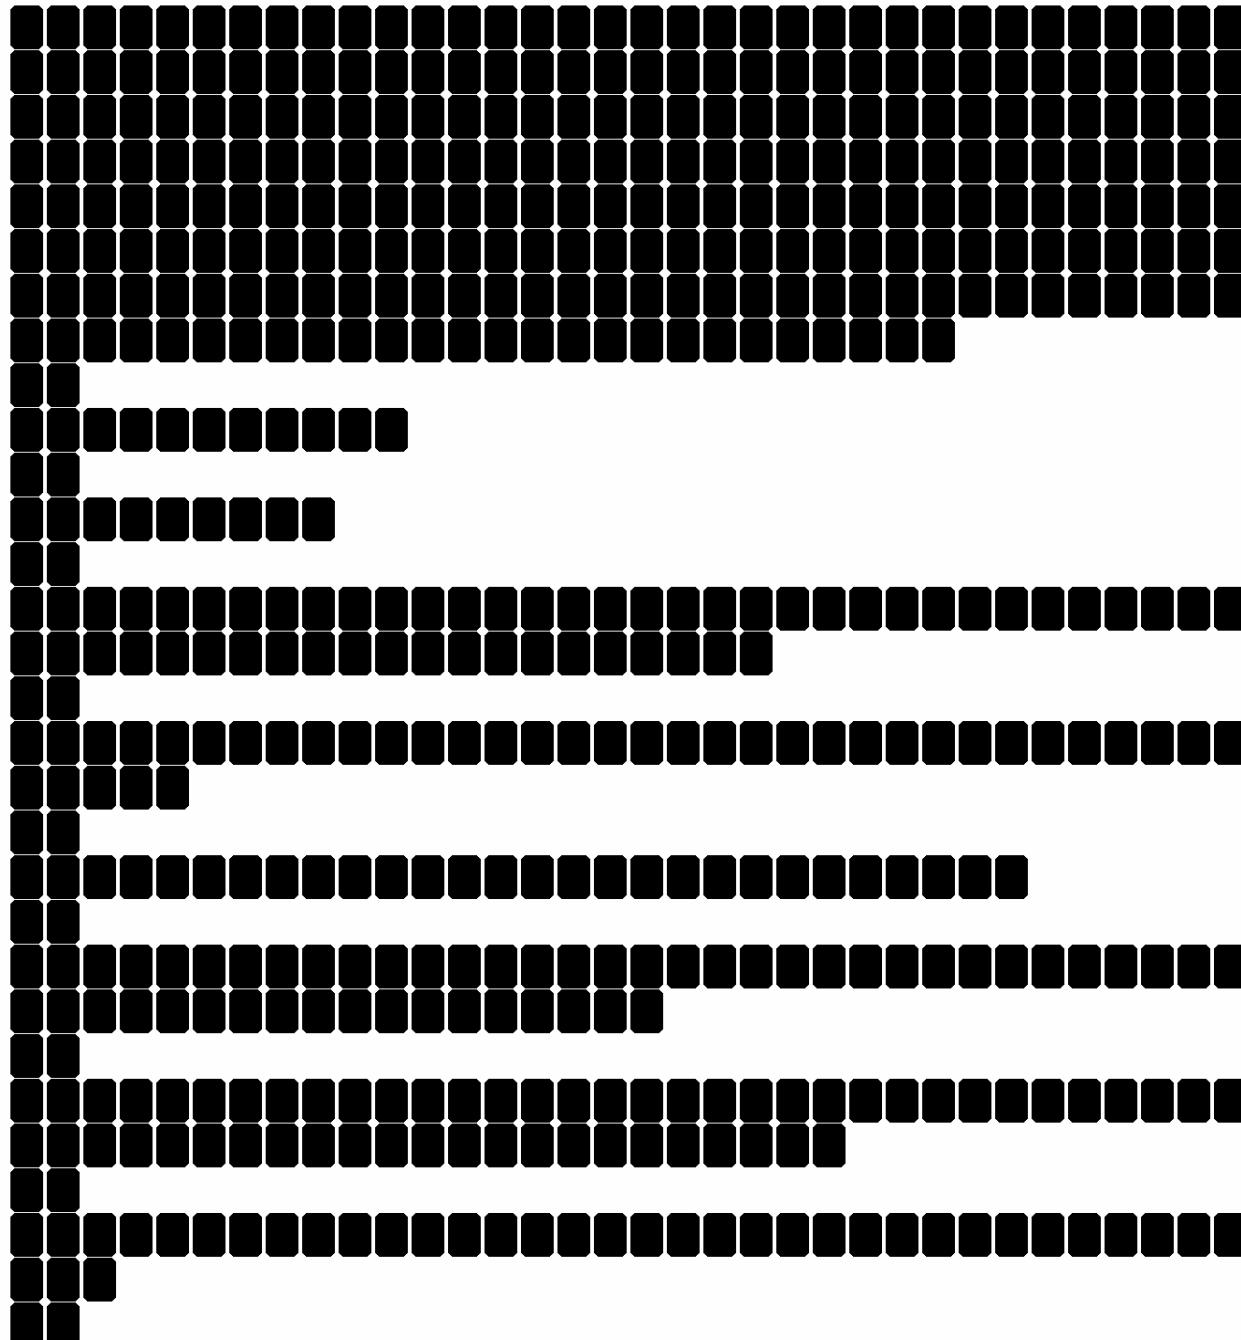
Janet's ducks lay **16** eggs per day. She eats **three** for breakfast every morning and bakes muffins for her friends everyday with **four**. She sells the remainder at the farmers' market daily for **\$2** per fresh duck egg. How much in dollars does she make every day at the farmers' market?

During
Supervised Fine
Tuning of LLM

During LLM
Inference



Step 0 / 4097
Tokens visible: 0 / 1024
Diffusion time \approx 1.0000
Frame 1 of 180



Paper:

openreview.net/pdf?id=E8adS5srd



Project Page:

anchored-diffusion-llm.github.io



Source Code:

github.com/LituRout/ADLM